

Towards deeper understandings of deep learning

Yuanzhi Li

Stanford University

date: Today

A bit of history of machine learning

• → 1940s

- 13.772 billion years BC to 1940s:

A bit of history of machine learning

• → 1940s

- 13.772 billion years BC to 1940s:
 - Nothing happened...

A bit of history of machine learning

• → 1940s

- 13.772 billion years BC to 1940s:
 - Nothing happened...
- 1950s to 2010s:

A bit of history of machine learning

• → 1940s

- 13.772 billion years BC to 1940s:
 - Nothing happened...
- 1950s to 2010s:
 - The age of **perceptrons**.

A bit of history of machine learning

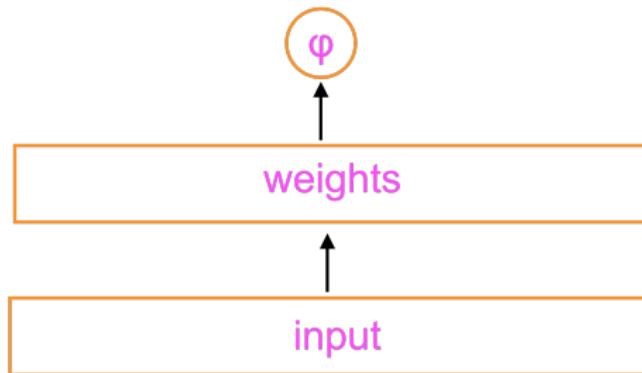
• → 1940s

- 13.772 billion years BC to 1940s:
 - Nothing happened...
- 1950s to 2010s:
 - The age of **perceptrons**.
 - Single layer learning models:

A bit of history of machine learning

• → 1940s

- 13.772 billion years BC to 1940s:
 - Nothing happened...
- 1950s to 2010s:
 - The age of **perceptrons**.
 - Single layer learning models:



A bit of history of machine learning

- Many variants of perceptrons: SVM, kernel SVM, products of experts

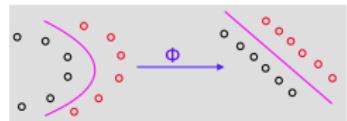
A bit of history of machine learning

- Many variants of perceptrons: SVM, kernel SVM, products of experts
- There is a lot of theory for them:

A bit of history of machine learning

- Many variants of perceptrons: SVM, kernel SVM, products of experts
- There is a lot of theory for them:

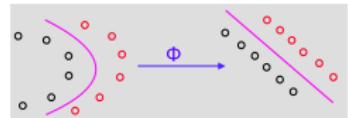
- Theory of kernels (representation power).



A bit of history of machine learning

- Many variants of perceptrons: SVM, kernel SVM, products of experts
- There is a lot of theory for them:

- Theory of kernels (representation power).



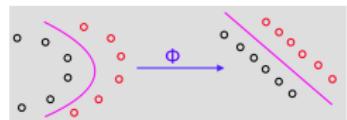
- Convex optimization (training time bounds).




A bit of history of machine learning

- Many variants of perceptrons: SVM, kernel SVM, products of experts
- There is a lot of **theory** for them:

- Theory of kernels (representation power).



- Convex optimization (training time bounds).
- VC theory (sample complexity and generalization).



A bit of history of machine learning

- In the past few years:

A bit of history of machine learning

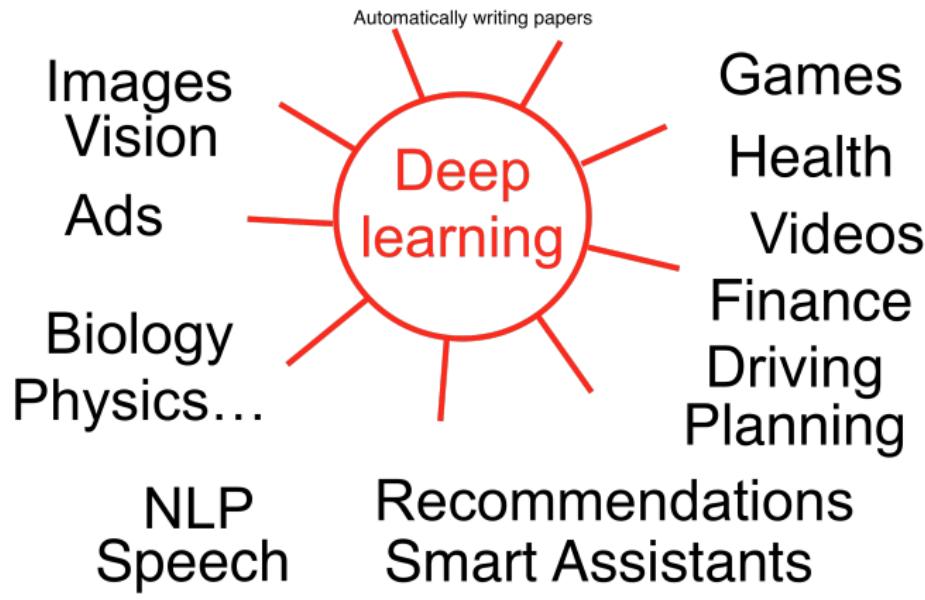
- In the past few years:
- Neural networks (Deep learning).

A bit of history of machine learning

- In the past few years:
- Neural networks (Deep learning).
- Works extremely well in **practice**, tons of **applications**.

A bit of history of machine learning

- In the past few years:
- Neural networks (Deep learning).
- Works extremely well in **practice**, tons of **applications**.



Neural Networks

- Multi-layer models.

Neural Networks

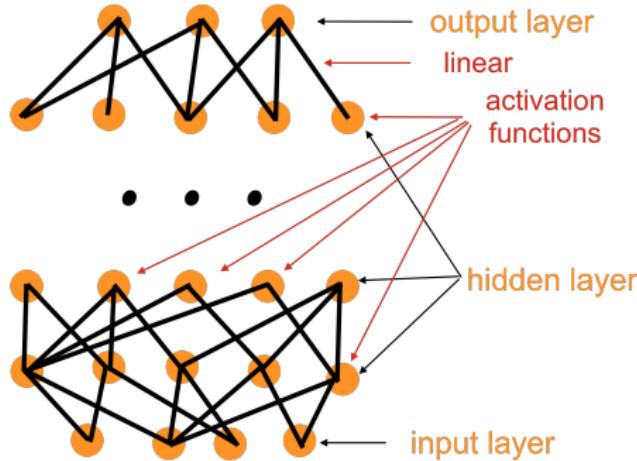
- Multi-layer models.
 - Input layer, output layer, and hidden layers.

Neural Networks

- Multi-layer models.
 - Input layer, output layer, and hidden layers.
 - Each hidden layer is associated with activation function(s).

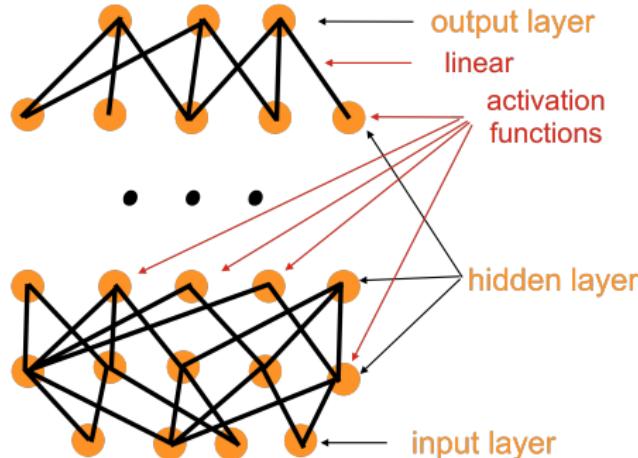
Neural Networks

- Multi-layer models.
 - Input layer, output layer, and hidden layers.
 - Each hidden layer is associated with activation function(s).



Neural Networks

- Multi-layer models.
 - Input layer, output layer, and hidden layers.
 - Each hidden layer is associated with activation function(s).



- In this talk, I will particularly focus on neural networks with **ReLU** activations.

ReLU

Neural Networks work extremely well in practice

- But how do I **prove** that my neural networks work? How do we **reason** about their powers?

Neural Networks work extremely well in practice

- But how do I **prove** that my neural networks work? How do we **reason** about their powers?
- In the past years, people have reached a satisfying answer to this question:

Neural Networks work extremely well in practice

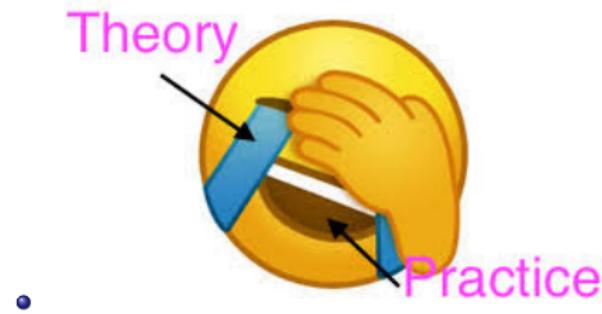
- But how do I **prove** that my neural networks work? How do we **reason** about their powers?
- In the past years, people have reached a satisfying answer to this question:
 - Stop asking about it!

Neural Networks work extremely well in practice

- But how do I **prove** that my neural networks work? How do we **reason** about their powers?
- In the past years, people have reached a satisfying answer to this question:
 - Stop asking about it!
 - Fact: Neural networks work in practice although we **can not fully prove it.**

Neural Networks work extremely well in practice

- But how do I **prove** that my neural networks work? How do we **reason** about their powers?
 - In the past years, people have reached a satisfying answer to this question:
 - Stop asking about it!
 - Fact: Neural networks work in practice although we **can not fully prove it.**



Why stop asking?

- There is a simple reason we might want to stop asking about it:

Why stop asking?

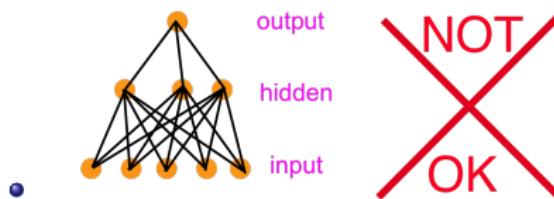
- There is a simple reason we might want to stop asking about it:
- In theory, neural **networks** are also well known as neural **notworks**.

Why stop asking?

- There is a simple reason we might want to stop asking about it:
- In theory, neural **networks** are also well known as neural **notworks**.
 - Learning a neural network with **three hidden units** is NP-hard (BR'98).


Why stop asking?

- There is a simple reason we might want to stop asking about it:
- In theory, neural **networks** are also well known as neural **notworks**.
 - Learning a neural network with **three hidden units** is NP-hard (BR'98).



Why stop asking?

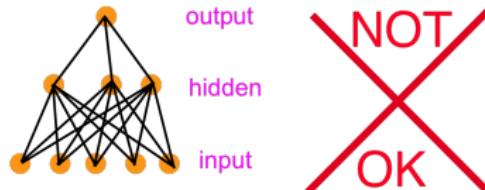
- There is a simple reason we might want to stop asking about it:
- In theory, neural **networks** are also well known as neural **notworks**.
 - Learning a neural network with **three hidden units** is NP-hard (BR'98).




- In practice, we have millions and they still **work**.

Why stop asking?

- There is a simple reason we might want to stop asking about it:
- In theory, neural **networks** are also well known as neural **notworks**.
 - Learning a neural network with **three hidden units** is NP-hard (BR'98).



- In practice, we have millions and they still **work**.



And even harder

- Yet there is another reason:

And even harder

- Yet there is another reason:
- In neural networks, we optimize any neural networks by essentially the same **simple simple simple simple simple** algorithm: **SGD**.

And even harder

- Yet there is another reason:
- In neural networks, we optimize any neural networks by essentially the same **simple simple simple simple simple** algorithm: **SGD**.
 - Starting from W_0 at random (often Gaussian).

And even harder

- Yet there is another reason:
- In neural networks, we optimize any neural networks by essentially the same **simple simple simple simple simple** algorithm: **SGD**.
 - Starting from W_0 at random (often Gaussian).
 - In each iteration, sample an example x from the data set, then update the weights W_t as: $W_{t+1} = W_t - \eta \nabla f(W_t, x)$.

And even harder

- Yet there is another reason:
- In neural networks, we optimize any neural networks by essentially the same **simple simple simple simple simple** algorithm: **SGD**.
 - Starting from W_0 at random (often Gaussian).
 - In each iteration, sample an example x from the data set, then update the weights W_t as: $W_{t+1} = W_t - \eta \nabla f(W_t, x)$.
- We use it (or its variants) almost everywhere, **regardless** of the problem/difficulty.

And even harder

- Yet there is another reason:
 - In neural networks, we optimize any neural networks by essentially the same **simple simple simple simple simple** algorithm: **SGD**.
 - Starting from W_0 at random (often Gaussian).
 - In each iteration, sample an example x from the data set, then update the weights W_t as: $W_{t+1} = W_t - \eta \nabla f(W_t, x)$.
 - We use it (or its variants) almost everywhere, **regardless** of the problem/difficulty.

```
optimizer = optim.SGD(para, lr)
optimizer.zero_grad()
loss.backward()
optimizer.step()
```

```
optimizer = optim.SGD(para, lr)
optimizer.zero_grad()
loss.backward()
optimizer.step()
```

Face recognition

And even harder

- Yet there is another reason:
- In neural networks, we optimize any neural networks by essentially the same simple simple simple simple simple algorithm: SGD.
 - Starting from W_0 at random (often Gaussian).
 - In each iteration, sample an example x from the data set, then update the weights W_t as: $W_{t+1} = W_t - \eta \nabla f(W_t, x)$.
- We use it (or its variants) almost everywhere, regardless of the problem/difficulty.

Nope!

```
optimizer = optim.SGD(para, lr)
optimizer.zero_grad()
loss.backward()
optimizer.step()
```

Yes!

Face
recognition

- Makes the theoretical work even harder (understand the universality of SGD).

My work

- In Machine learning, many models are NP-hard, but nevertheless solved **efficiently** in practice by **simple algorithms**.

Graphical Models Tensor decomposition
k-means Mixture of regressions
ICA NMF
Deep Learning
Bellman equation Sparse Coding Weighted SVD
Topic Modeling MLE Dictionary Matrix Completion
 Learning Matrix sensing

My work

- In Machine learning, many models are NP-hard, but nevertheless solved **efficiently** in practice by **simple algorithms**.

Graphical Models Tensor decomposition
k-means Mixture of regressions
ICA NMF Deep Learning
Bellman equation Sparse Coding Weighted SVD
Topic Modeling MLE Dictionary Matrix Completion
 Learning Matrix sensing

- I work on new, **principled ways** to understand and enhance the success of deep learning/machine learning in general.

My work

- In Machine learning, many models are NP-hard, but nevertheless solved **efficiently** in practice by **simple algorithms**.



- I work on new, **principled ways** to understand and enhance the success of deep learning/machine learning in general.
- In this talk I will focus on principles relevant to (multi-layer) neural networks with **ReLU** activation functions (NP-hard model).

My work

- In Machine learning, many models are NP-hard, but nevertheless solved **efficiently** in practice by **simple algorithms**.



- I work on new, **principled ways** to understand and enhance the success of deep learning/machine learning in general.
- In this talk I will focus on principles relevant to (multi-layer) neural networks with **ReLU** activation functions (NP-hard model).
- Training: **SGD** with random initializations (**simple algorithm**).

My work

- In Machine learning, many models are NP-hard, but nevertheless solved **efficiently** in practice by **simple algorithms**.



- I work on new, **principled ways** to understand and enhance the success of deep learning/machine learning in general.
- In this talk I will focus on principles relevant to (multi-layer) neural networks with **ReLU** activation functions (NP-hard model).
- Training: **SGD** with random initializations (**simple algorithm**).
- We will discuss two principles in this talk.

My work

- In Machine learning, many models are NP-hard, but nevertheless solved **efficiently** in practice by **simple algorithms**.



- I work on new, **principled ways** to understand and enhance the success of deep learning/machine learning in general.
- In this talk I will focus on principles relevant to (multi-layer) neural networks with **ReLU** activation functions (NP-hard model).
- Training: **SGD** with random initializations (**simple algorithm**).
- We will discuss two principles in this talk.



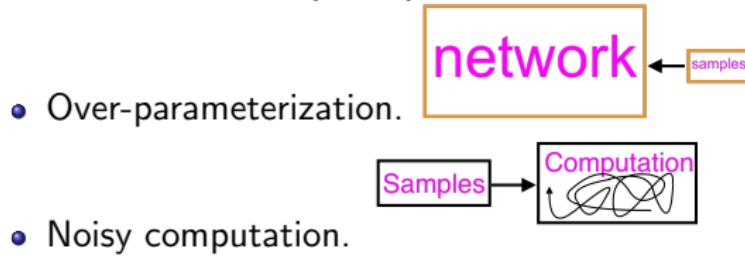
- Over-parameterization.

My work

- In Machine learning, many models are NP-hard, but nevertheless solved **efficiently** in practice by **simple algorithms**.



- I work on new, **principled ways** to understand and enhance the success of deep learning/machine learning in general.
- In this talk I will focus on principles relevant to (multi-layer) neural networks with **ReLU** activation functions (NP-hard model).
- Training: **SGD** with random initializations (**simple algorithm**).
- We will discuss two principles in this talk.



My work

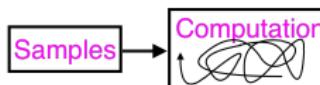
- In Machine learning, many models are NP-hard, but nevertheless solved **efficiently** in practice by **simple algorithms**.



- I work on new, **principled ways** to understand and enhance the success of deep learning/machine learning in general.
- In this talk I will focus on principles relevant to (multi-layer) neural networks with **ReLU** activation functions (NP-hard model).
- Training: **SGD** with random initializations (**simple algorithm**).
- We will discuss two principles in this talk.



- Over-parameterization.



- Noisy computation.
- And how can they help in learning neural networks, **provably**.

Conversation between deep learning and theory

- (1). My network finds the global optimal on the training set, but it converges slowly. What should I do?

Conversation between deep learning and theory

- (1). My network finds the global optimal on the training set, but it converges slowly. What should I do?
 - (A person who just took the convex optimization course): Make the total number of parameters **smaller** so it runs faster.

Conversation between deep learning and theory

- (1). My network finds the global optimal on the training set, but it converges slowly. What should I do?
 - (A person who just took the convex optimization course): Make the total number of parameters **smaller** so it runs faster.
 - (A person who works on deep learning): You should use a **larger** network.

Conversation between deep learning and theory

- (1). My network finds the global optimal on the training set, but it converges slowly. What should I do?
 - (A person who just took the convex optimization course): Make the total number of parameters **smaller** so it runs faster.
 - (A person who works on deep learning): You should use a **larger** network.
- (2). My network finds the global optimal on the training set, but it generalizes badly. What should I do?

Conversation between deep learning and theory

- (1). My network finds the global optimal on the training set, but it converges slowly. What should I do?
 - (A person who just took the convex optimization course): Make the total number of parameters **smaller** so it runs faster.
 - (A person who works on deep learning): You should use a **larger** network.
- (2). My network finds the global optimal on the training set, but it generalizes badly. What should I do?
 - (A person who just learned the VC dimension): Make the total number of parameters **smaller** so it generalizes better.

Conversation between deep learning and theory

- (1). My network finds the global optimal on the training set, but it converges slowly. What should I do?
 - (A person who just took the convex optimization course): Make the total number of parameters **smaller** so it runs faster.
 - (A person who works on deep learning): You should use a **larger** network.
- (2). My network finds the global optimal on the training set, but it generalizes badly. What should I do?
 - (A person who just learned the VC dimension): Make the total number of parameters **smaller** so it generalizes better.
 - (A person who works on deep learning): You should use a **larger** network.

Conversation between deep learning and theory

- (1). My network finds the global optimal on the training set, but it converges slowly. What should I do?
 - (A person who just took the convex optimization course): Make the total number of parameters **smaller** so it runs faster.
 - (A person who works on deep learning): You should use a **larger** network.
- (2). My network finds the global optimal on the training set, but it generalizes badly. What should I do?
 - (A person who just learned the VC dimension): Make the total number of parameters **smaller** so it generalizes better.
 - (A person who works on deep learning): You should use a **larger** network.
- (3). My network finds a good solution empirically, but I can't prove it. What should I do?

Conversation between deep learning and theory

- (1). My network finds the global optimal on the training set, but it converges slowly. What should I do?
 - (A person who just took the convex optimization course): Make the total number of parameters **smaller** so it runs faster.
 - (A person who works on deep learning): You should use a **larger** network.
- (2). My network finds the global optimal on the training set, but it generalizes badly. What should I do?
 - (A person who just learned the VC dimension): Make the total number of parameters **smaller** so it generalizes better.
 - (A person who works on deep learning): You should use a **larger** network.
- (3). My network finds a good solution empirically, but I can't prove it. What should I do?
 - (In this talk): You should use a **larger** network!

Principle of over-parameterization

- By building up a network with **(much more) parameters** than the total number of training examples.

Principle of over-parameterization

- By building up a network with (much more) parameters than the total number of training examples.



Principle of over-parameterization

- By building up a network with (much more) parameters than the total number of training examples.



- Improves both the training and generalization.

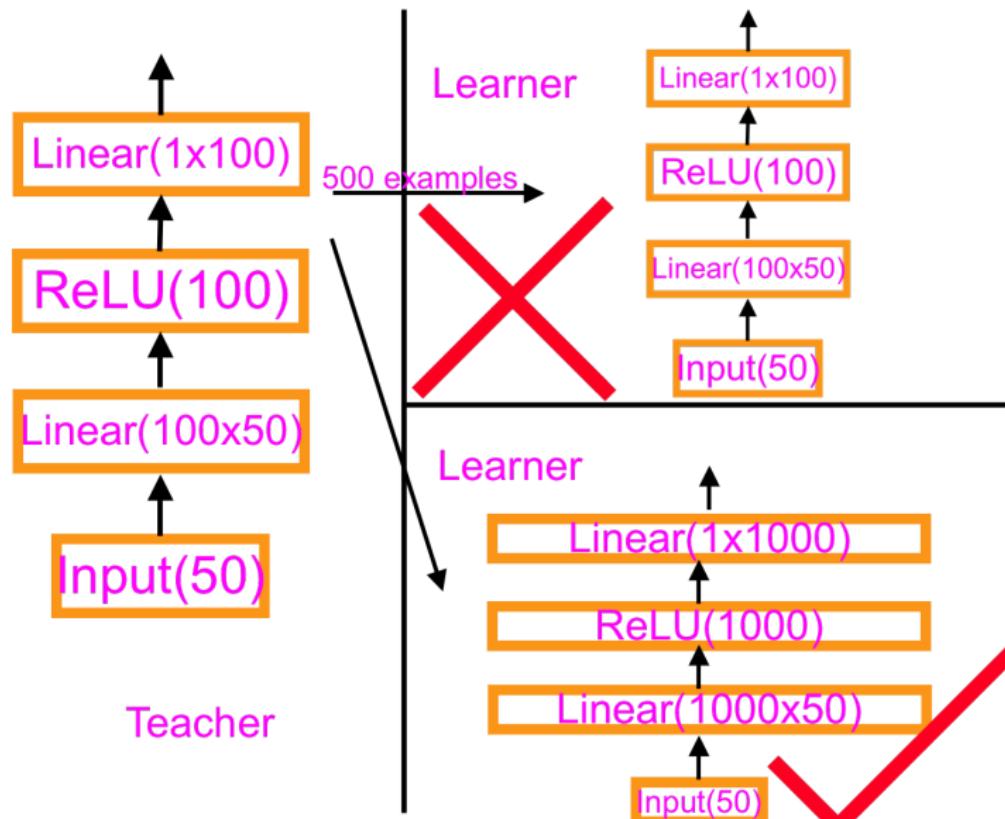
Principle of over-parameterization

- By building up a network with (much more) parameters than the total number of training examples.



- Improves both the training and generalization.
- And it improves the theory.

Folklore example for training.



Example for generalization

Widen factor	Number of parameters	Test error
1	0.6M	6.85
2	2.2M	5.33
4	8.9M	4.97
8	36.5M	4.66

Table: Depth 40 WideResNet on CIFAR-10 (0.05M training examples), training errors are all 0.

Example of the Theorem

- Given enough over-parameterization, prove that:

Example of the Theorem

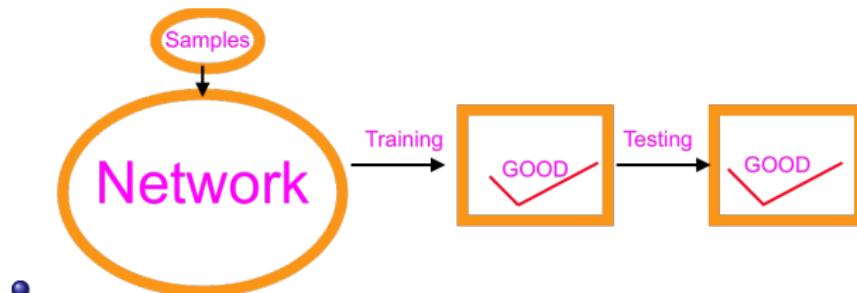
- Given enough over-parameterization, prove that:
- (1). SGD will find a good solution on the training data set (close to zero training error).

Example of the Theorem

- Given enough over-parameterization, prove that:
- (1). SGD will find a good solution on the training data set (close to zero training error).
- (2). And it generalizes to test data set.


Example of the Theorem

- Given enough over-parameterization, prove that:
- (1). SGD will find a good solution on the training data set (close to zero training error).
- (2). And it generalizes to test data set.



Example of the Theorem

- Given enough over-parameterization, prove that:
- (1). SGD will find a good solution on the training data set (close to zero training error).
- (2). And it generalizes to test data set.



- We begin with our theorem for (1), then we will see (2) as well.

Our Theorem

Theorem (Sketched, (LL'18, ALS'18a,b))

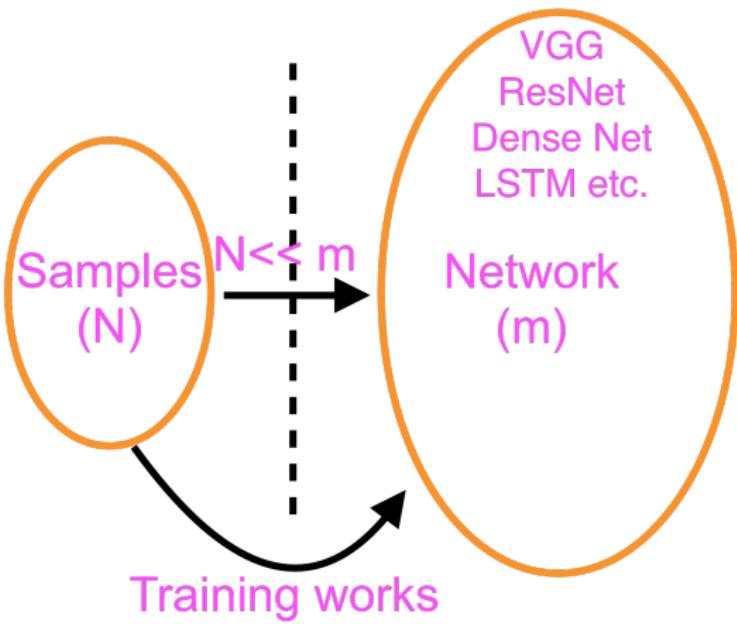
Given N different training examples x_1, \dots, x_N with labels y_1, \dots, y_N , then for every $\varepsilon > 0$, as long as the number of neurons (m) in the network satisfies

$$m \geq \text{poly}(N \log(1/\varepsilon))$$

then SGD starting from gaussian random initialization finds an ε -approximate global optimal of the training objective in time $\text{poly}(m/\varepsilon)$.

The theorem holds for multi-layer DNNs, CNNs, ResNet and Recurrent Neural Networks (all with ReLU activation functions), the training loss can be given by ℓ_2 loss, cross entropy etc.

Our theorem by picture



Closer look at the Theorem:

- $m \geq \text{poly}(N \log(1/\varepsilon))$ implies good fitting on the training data set.

Closer look at the Theorem:

- $m \geq \text{poly}(N \log(1/\varepsilon))$ implies good fitting on the training data set.
- Essentially no assumptions on the training examples, in particular training labels can be **random**.

Closer look at the Theorem:

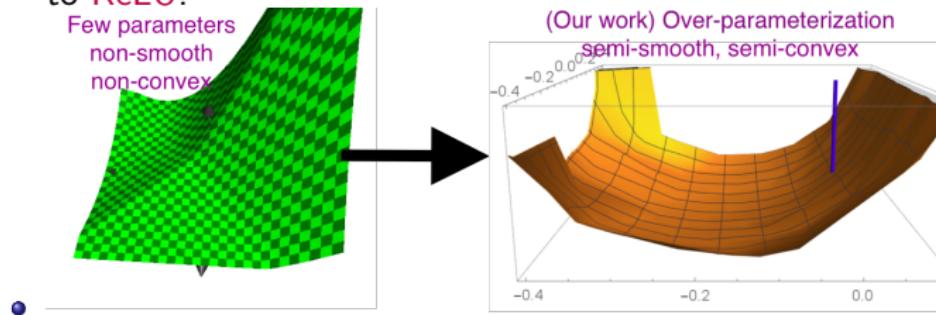
- $m \geq \text{poly}(N \log(1/\varepsilon))$ implies good fitting on the training data set.
- Essentially no assumptions on the training examples, in particular training labels can be **random**.
 - From capacity perspective: Obviously there **exists** a network that fits the training data.

Closer look at the Theorem:

- $m \geq \text{poly}(N \log(1/\varepsilon))$ implies good fitting on the training data set.
- Essentially no assumptions on the training examples, in particular training labels can be **random**.
 - From capacity perspective: Obviously there **exists** a network that fits the training data.
 - From optimization perspective: How can **SGD** find such fitting? The training objective is not only **non-convex**, but **non-smooth** as well due to **ReLU**.

Closer look at the Theorem:

- $m \geq \text{poly}(N \log(1/\varepsilon))$ implies good fitting on the training data set.
- Essentially no assumptions on the training examples, in particular training labels can be **random**.
 - From capacity perspective: Obviously there **exists** a network that fits the training data.
 - From optimization perspective: How can **SGD** find such fitting? The training objective is not only **non-convex**, but **non-smooth** as well due to **ReLU**.



Training is ok with over-parameterization

- Memorization is not a completely trivial theorem, but much simpler compare to this question:

Training is ok with over-parameterization

- Memorization is not a completely trivial theorem, but much simpler compare to this question:
 - What about generalization?

Training is ok with over-parameterization

- Memorization is not a completely trivial theorem, but much simpler compare to this question:
 - What about generalization?
- The empirical example (ZBHR'16), that is also proved by our theorem.

Training is ok with over-parameterization

- Memorization is not a completely trivial theorem, but much simpler compare to this question:
 - What about generalization?
- The empirical example (ZBHR'16), that is also proved by our theorem.
 - Over-parameterized networks (AlexNet) can fit CIFAR-10 data with **random labels**.

Training is ok with over-parameterization

- Memorization is not a completely trivial theorem, but much simpler compare to this question:
 - What about generalization?
- The empirical example (ZBHR'16), that is also proved by our theorem.
 - Over-parameterized networks (AlexNet) can fit CIFAR-10 data with **random labels**.
 - The **capacity** of the model is **way larger** than the total number of training examples.

Training is ok with over-parameterization

- Memorization is not a completely trivial theorem, but much simpler compare to this question:
 - What about generalization?
- The empirical example (ZBHR'16), that is also proved by our theorem.
 - Over-parameterized networks (AlexNet) can fit CIFAR-10 data with **random labels**.
 - The **capacity** of the model is **way larger** than the total number of training examples.
 - More importantly, **SGD can find** such (over)fitting.

Training is ok with over-parameterization

- Memorization is not a completely trivial theorem, but much simpler compare to this question:
 - What about generalization?
- The empirical example (ZBHR'16), that is also proved by our theorem.
 - Over-parameterized networks (AlexNet) can fit CIFAR-10 data with **random labels**.
 - The **capacity** of the model is **way larger** than the total number of training examples.
 - More importantly, **SGD can find** such (over)fitting.



•

Training is ok with over-parameterization

- Memorization is not a completely trivial theorem, but much simpler compare to this question:
 - What about generalization?
- The empirical example (ZBHR'16), that is also proved by our theorem.
 - Over-parameterized networks (AlexNet) can fit CIFAR-10 data with **random labels**.
 - The **capacity** of the model is **way larger** than the total number of training examples.
 - More importantly, **SGD can find** such (over)fitting.



- So why does the solution found by **SGD** still generalize?

Two main reasons for generalization:

- The labels are not random (necessary).

Two main reasons for generalization:

- The labels are not random (necessary).
 - What about SGD?

Two main reasons for generalization:

- The labels are not random (necessary).
 - What about SGD?
- The **inductive bias** of SGD: SGD usually biases towards **generalizable** solutions instead of the solutions that simply memorize the training data.

Two main reasons for generalization:

- The labels are not random (necessary).
 - What about SGD?
- The **inductive bias** of SGD: SGD usually biases towards **generalizable** solutions instead of the solutions that simply memorize the training data.
 - We are going to prove these claims for certain neural networks.

Two main reasons for generalization:


- The labels are not random (necessary).
 - What about SGD?
- The **inductive bias** of SGD: SGD usually biases towards **generalizable** solutions instead of the solutions that simply memorize the training data.
 - We are going to prove these claims for certain neural networks.
 - We start with two layer networks, and then multi-layer ones.

Labels are not random? (Two layer networks)

- Assume that the labels are realizable by a simple neural network:


Labels are not random? (Two layer networks)

- Assume that the labels are realizable by a simple neural network:



Labels are not random? (Two layer networks)

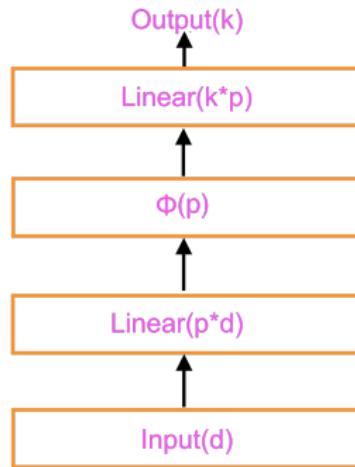
- Assume that the labels are realizable by a simple neural network:



- Two-layer network with p hidden neurons, where p is something small (e.g. 100). Φ is a smooth activation function (such as \sin , \cos , \exp , \tanh , *sigmoid*, and low degree polynomials).

Labels are not random? (Two layer networks)

- Assume that the labels are realizable by a simple neural network:



- Two-layer network with p hidden neurons, where p is something small (e.g. 100). Φ is a smooth activation function (such as \sin , \cos , \exp , \tanh , *sigmoid*, and low degree polynomials).
- Such that the average loss of this network on the **training data set** is $\leq \varepsilon$.

Our Theorem

Theorem (Sketched (LL'18, ALL'18))

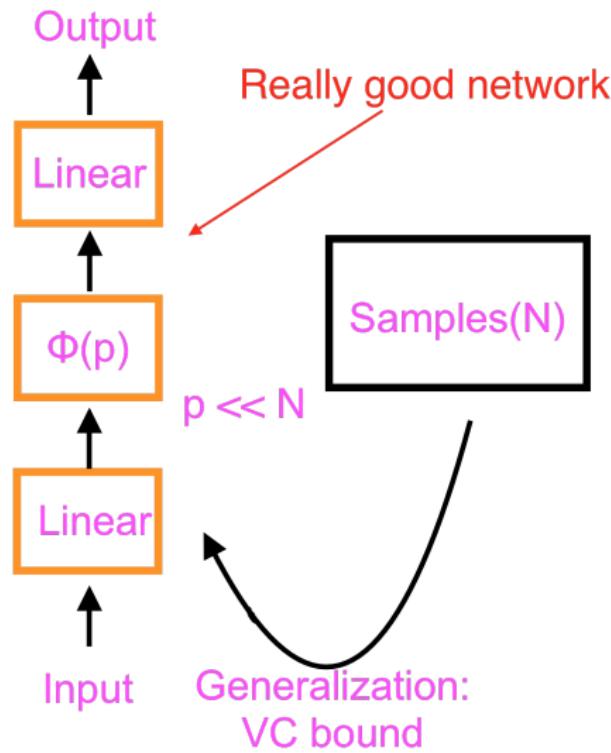
Suppose the labels of the data can be realized as in the previous slides, then as long as the number of training examples N satisfies:

$$N \geq \text{poly}(kp/\varepsilon) \times \text{poly}(\log(m))$$


Then **SGD** on a two layer neural network (with m neurons and **ReLU** activation functions) finds a solution with **generalization gap** $\leq \varepsilon$ in time $\text{poly}(m/\varepsilon)$.

Generalization gap = average error on test data - average error on training data.

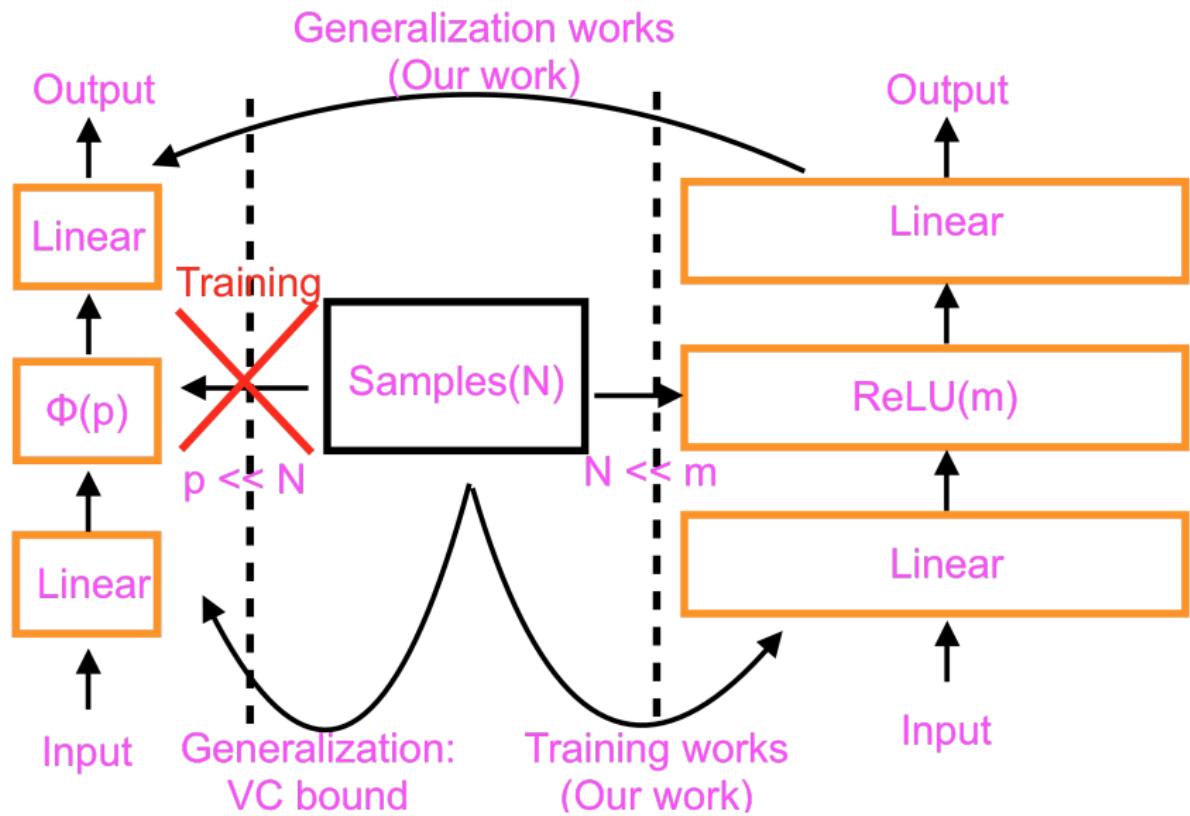
Picture of the Theorem



Picture of the Theorem



Picture of the Theorem



The main message

- Training a two layer network is NP-hard.

The main message

- Training a two layer network is NP-hard.
- But we are not optimizing a network up to optimal, we just want it to do **as well as** the best **(much)** smaller networks.


The main message

- Training a two layer network is NP-hard.
- But we are not optimizing a network up to optimal, we just want it to do **as well as the best (much) smaller networks**.
- We are training a network with **$1M$** parameters to match the performance of a network with **$1K$** parameters, and that is **easy!**

The main message

- Training a two layer network is NP-hard.
- But we are not optimizing a network up to optimal, we just want it to do **as well as the best (much) smaller networks**.
- We are training a network with **$1M$** parameters to match the performance of a network with **$1K$** parameters, and that is **easy!**
- This is what behind the principle of over-parameterization.

Intuition



Generalization

$y = \cos(x)$

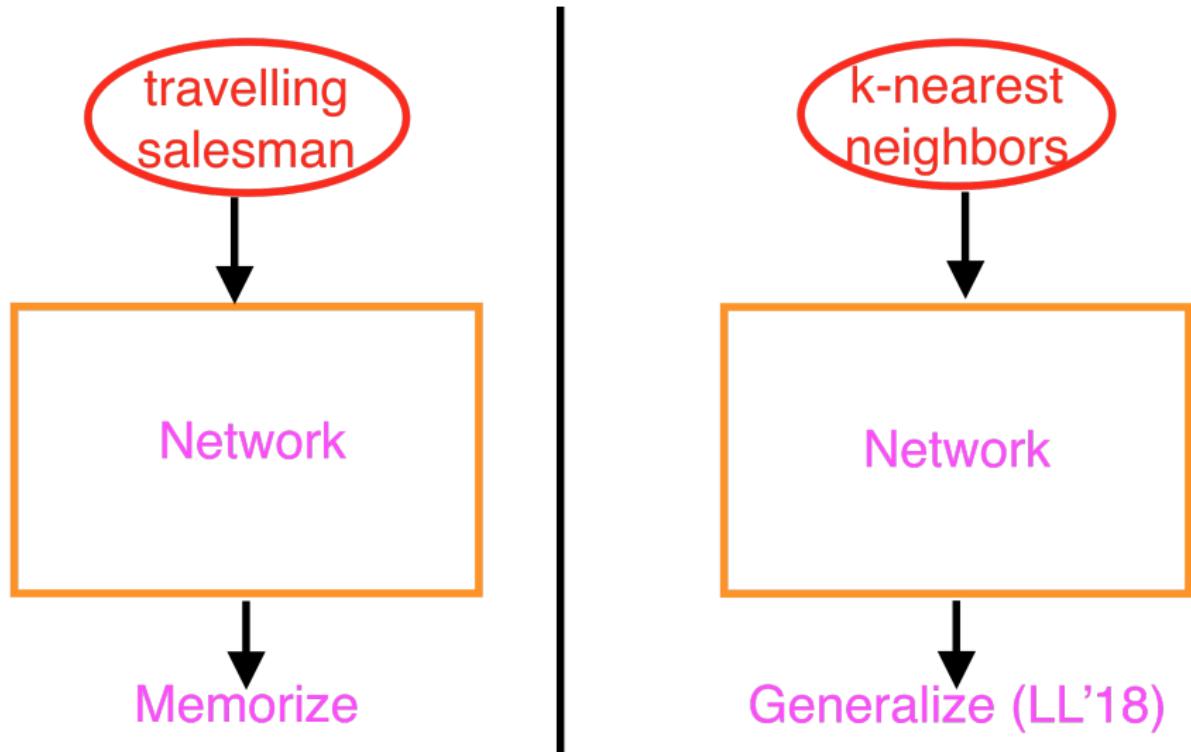
Simple!

Memorization

if $x = 1$, $y = 0.54$
if $x = 2$, $y = -0.42$
if $x = 3$, $y = -0.99$
if $x = 4$, $y = -0.65$
if $x = 5$, $y = 0.28$
if $x = 6$, $y = 0.96$
if $x = 7$, $y = 0.75$
if $x = 8$, $y = -0.15$
if $x = 9$, $y = -0.91$
if $x = 10$, $y = -0.84$

I gave up typing
those things...
too tired...

Intuition



Existence of small networks is somewhat necessary

Networks with more hidden layers?


- Our theorem [ALL'18] also works for three layer networks (two hidden layers), and ([ALL'18] Corollary) single skip ResNet with **multiple layers**.

Networks with more hidden layers?

- Our theorem [ALL'18] also works for three layer networks (two hidden layers), and ([ALL'18] Corollary) single skip ResNet with **multiple layers**.
- Theoretical reasoning on multi-layer networks is **much harder** due to the **non-convex interactions** between the hidden layers.

Networks with more hidden layers?

- Our theorem [ALL'18] also works for three layer networks (two hidden layers), and ([ALL'18] Corollary) single skip ResNet with **multiple layers**.
- Theoretical reasoning on multi-layer networks is **much harder** due to the **non-convex interactions** between the hidden layers.



Key Messages

- Depth **provably** matters.

Key Messages

- Depth **provably** matters.
- For the set of functions (super set of three layer networks with smooth activations) **efficiently learnable** by (over-paramterized) **three layer networks**.

Key Messages

- Depth **provably** matters.
- For the set of functions (super set of three layer networks with smooth activations) **efficiently learnable** by (over-paramterized) **three layer networks**.
 - The **ONLY** efficient learning algorithm we know is to train a (over-paramterized) three layer networks.

Key Messages

- Depth **provably** matters.
- For the set of functions (super set of three layer networks with smooth activations) **efficiently learnable** by (over-paramterized) **three layer networks**.
 - The **ONLY** efficient learning algorithm we know is to train a (over-paramterized) three layer networks.
 - Can not be learnt by kernel methods efficiently.


Key Messages

- Depth **provably** matters.
- For the set of functions (super set of three layer networks with smooth activations) **efficiently learnable** by (over-paramterized) **three layer networks**.
 - The **ONLY** efficient learning algorithm we know is to train a (over-paramterized) three layer networks.
 - Can not be learnt by kernel methods efficiently.
- Prior works:

Key Messages

- Depth **provably** matters.
- For the set of functions (super set of three layer networks with smooth activations) **efficiently learnable** by (over-paramterized) **three layer networks**.
 - The **ONLY** efficient learning algorithm we know is to train a (over-paramterized) three layer networks.
 - Can not be learnt by kernel methods efficiently.
- Prior works:

- Linear networks.




Key Messages

- Depth **provably** matters.
- For the set of functions (super set of three layer networks with smooth activations) **efficiently learnable** by (over-paramterized) **three layer networks**.
 - The **ONLY** efficient learning algorithm we know is to train a (over-paramterized) three layer networks.
 - Can not be learnt by kernel methods efficiently.
- Prior works:
 - Linear networks.
 - Networks on Gaussian inputs (tensor decomposition).

Key Messages

- Depth **provably** matters.
- For the set of functions (super set of three layer networks with smooth activations) **efficiently learnable** by (over-paramterized) **three layer networks**.
 - The **ONLY** efficient learning algorithm we know is to train a (over-paramterized) three layer networks.
 - Can not be learnt by kernel methods efficiently.
- Prior works:

- Linear networks.



- Networks on Gaussian inputs (tensor decomposition).




- Networks to learn linear separable data.



Key Messages

- Depth **provably** matters.
- For the set of functions (super set of three layer networks with smooth activations) **efficiently learnable** by (over-paramterized) **three layer networks**.
 - The **ONLY** efficient learning algorithm we know is to train a (over-paramterized) three layer networks.
 - Can not be learnt by kernel methods efficiently.
- Prior works:




- Linear networks.



- Networks on Gaussian inputs (tensor decomposition).



- Networks to learn linear separable data.



kernel.

Summary

- We have shown that:

Summary

- We have shown that:
 - SGD provably optimizes over-parameterization deep neural networks with ReLU activations, on the training data set.

Summary

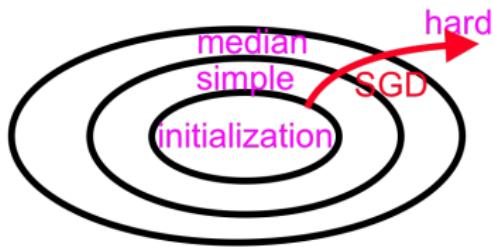
- We have shown that:
 - SGD provably optimizes over-parameterization deep neural networks with ReLU activations, on the training data set.
 - SGD also biases towards generalizable solutions on two/three layer networks and ResNet if the labels are structured.

Summary

- We have shown that:
 - SGD provably optimizes over-parameterization deep neural networks with ReLU activations, on the training data set.
 - SGD also biases towards generalizable solutions on two/three layer networks and ResNet if the labels are structured.
 - Reason: SGD biases towards simple solutions.

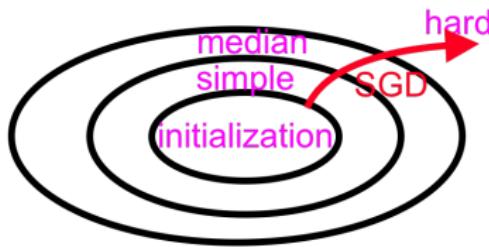
Summary

- We have shown that:
 - **SGD** provably optimizes over-parameterization deep neural networks with **ReLU** activations, on the training data set.
 - **SGD** also biases towards generalizable solutions on two/three layer networks and ResNet if the labels are structured.
 - Reason: **SGD** biases towards simple solutions.



Summary

- We have shown that:
 - SGD provably optimizes over-parameterization deep neural networks with ReLU activations, on the training data set.
 - SGD also biases towards generalizable solutions on two/three layer networks and ResNet if the labels are structured.
 - Reason: SGD biases towards simple solutions.



- Question: Does this bias always lead to best generalization? Can we do better?

SGD with learning rate decay

- When we train a neural network (VGG, ResNet, Dense Net, etc. for image classification etc.)

SGD with learning rate decay

- When we train a neural network (VGG, ResNet, Dense Net, etc. for image classification etc.)
- It is better to use large learning rate first, then decay the learning rate.

SGD with learning rate decay


- When we train a neural network (VGG, ResNet, Dense Net, etc. for image classification etc.)
- It is better to use large learning rate first, then decay the learning rate.



SGD with different learning rates


Job Name: cust-r-1-reg-reg-resnet-32-cifar10-sgd-sched81.122-mom0.9

Username	StartTime	Elapsed	VC-Queue	Build	ApplicationID	GPUs	Debug	Retries	Preempts	Status
zeyuana	2019-01-13 02:57:20	2:03:06	msrlabs-parallel		1545093298057_6029	1	false	0	0	Running



Name: cust-r-1-reg-reg-resnet-32-cifar10-sgd-lr0.001-epoch5000-sched5000.5001-mom0.9

Username	StartTime	Elapsed	VC-Queue	Build	ApplicationID	GPUs	Debug	Ref
zeyuana	2019-01-14 03:09:50	27:57:57	msrlabs-parallel		1545093298057_6453	1	false	



Principle of “Noisy computation”

- Training loss of using large learning rate + learning rate decay and using small learning rate are all close to zero.

Principle of “Noisy computation”


- Training loss of using large learning rate + learning rate decay and using small learning rate are all close to zero.
- But large learning rate generalizes better!

Principle of “Noisy computation”

- Training loss of using large learning rate + learning rate decay and using small learning rate are all close to zero.
- But large learning rate generalizes better!
- Principle: It is better to incorporate noise when training neural networks, as it improves generalization. ($SGD + \text{large learning rate} = \text{large noise}$)

Principle of “Noisy computation”


- Training loss of using large learning rate + learning rate decay and using small learning rate are all close to zero.
- But large learning rate generalizes better!
- Principle: It is better to incorporate noise when training neural networks, as it improves generalization. ($SGD + \text{large learning rate} = \text{large noise}$)



-

Principle of “Noisy computation”

- Training loss of using large learning rate + learning rate decay and using small learning rate are all close to zero.
- But large learning rate generalizes better!
- Principle: It is better to incorporate noise when training neural networks, as it improves generalization. ($SGD + \text{large learning rate} = \text{large noise}$)



-
- But why?

Our current work

- The principle of “noisy computation”:

Our current work

- The principle of “noisy computation”:
- On certain data sets, when training a neural network using SGD, large learning rate + learning rate decay provably generalizes better than small learning rate.

Example of the data set

- Texts labeled as happy:

Example of the data set

- Texts labeled as happy:
 - I am so happy ☺

Example of the data set

- Texts labeled as happy:
 - I am so happy 😊
 - Feeling good 😊

Example of the data set

- Texts labeled as happy:

- I am so happy ☺
- Feeling good ☺
- Excited ☺

Example of the data set

- Texts labeled as happy:
 - I am so happy 😊
 - Feeling good 😊
 - Excited 😊
 - Today is a good day!

Example of the data set

- Texts labeled as happy:
 - I am so happy 😊
 - Feeling good 😊
 - Excited 😊
 - Today is a good day!
 - I am Groot 😊

Example of the data set

- Texts labeled as happy:
 - I am so happy 😊
 - Feeling good 😊
 - Excited 😊
 - Today is a good day!
 - I am Groot 😊
 - ...

Example of the data set

- Texts labeled as happy:
 - I am so happy 😊
 - Feeling good 😊
 - Excited 😊
 - Today is a good day!
 - I am Groot 😊
 - ...
- Texts labeled as sad:

Example of the data set

- Texts labeled as happy:
 - I am so happy 😊
 - Feeling good 😊
 - Excited 😊
 - Today is a good day!
 - I am Groot 😊
 - ...
- Texts labeled as sad:
 - I am so unhappy 😞

Example of the data set

- Texts labeled as happy:
 - I am so happy 😊
 - Feeling good 😊
 - Excited 😊
 - Today is a good day!
 - I am Groot 😊
 - ...
- Texts labeled as sad:
 - I am so unhappy 😢
 - Feeling sorrowful 😢

Example of the data set

- Texts labeled as happy:
 - I am so happy 😊
 - Feeling good 😊
 - Excited 😊
 - Today is a good day!
 - I am Groot 😊
 - ...
- Texts labeled as sad:
 - I am so unhappy 😞
 - Feeling sorrowful 😞
 - Very unhappy.

Example of the data set

- Texts labeled as happy:
 - I am so happy 😊
 - Feeling good 😊
 - Excited 😊
 - Today is a good day!
 - I am Groot 😊
 - ...
- Texts labeled as sad:
 - I am so unhappy 😢
 - Feeling sorrowful 😢
 - Very unhappy.
 - Being wrecked 😢

Example of the data set

- Texts labeled as happy:
 - I am so happy 😊
 - Feeling good 😊
 - Excited 😊
 - Today is a good day!
 - I am Groot 😊
 - ...
- Texts labeled as sad:
 - I am so unhappy 😢
 - Feeling sorrowful 😢
 - Very unhappy.
 - Being wrecked 😢
 - I am Groot 😢

Example of the data set

- Texts labeled as happy:
 - I am so happy 😊
 - Feeling good 😊
 - Excited 😊
 - Today is a good day!
 - I am Groot 😊
 - ...
- Texts labeled as sad:
 - I am so unhappy 😢
 - Feeling sorrowful 😢
 - Very unhappy.
 - Being wrecked 😢
 - I am Groot 😢
 - ...

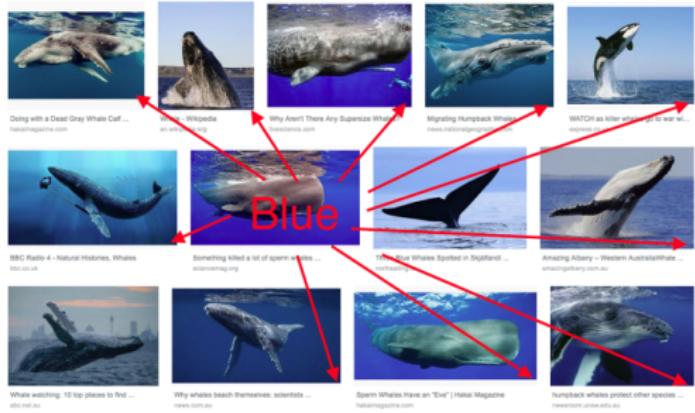
Example of the data set

- Texts labeled as happy:

- I am so happy 😊
- Feeling good 😊
- Excited 😊
- Today is a good day!
- I am Groot 😊
- ...

- Texts labeled as sad:

- I am so unhappy 😞
- Feeling sorrowful 😞
- Very unhappy.
- Being wrecked 😞
- I am Groot 😞
- ...



The intuition behind SGD

- What's wrong with SGD using small learning rate?

The intuition behind SGD


- What's wrong with **SGD** using small learning rate?
- **SGD** with small learning rate will quickly **memorize** the special symbols ☺, ☹.

The intuition behind SGD

- What's wrong with SGD using small learning rate?
- SGD with small learning rate will quickly memorize the special symbols 😊, 😆.
- SGD then uses very few examples without these symbols to learn the sentences.


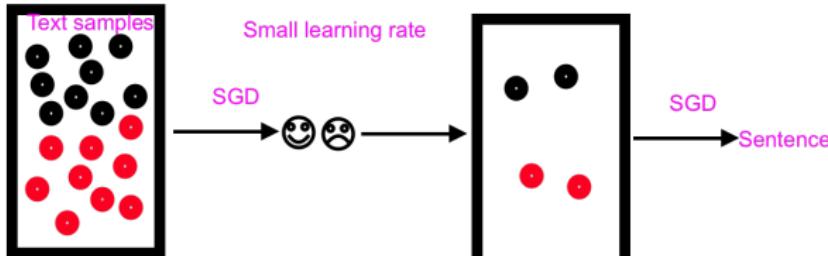
The intuition behind SGD

- What's wrong with SGD using small learning rate?
- SGD with small learning rate will quickly memorize the special symbols 😊, 😓.
- SGD then uses very few examples without these symbols to learn the sentences.



The intuition behind SGD

- What's wrong with SGD using small learning rate?
- SGD with small learning rate will quickly memorize the special symbols ☺, ☹.
- SGD then uses very few examples without these symbols to learn the sentences.

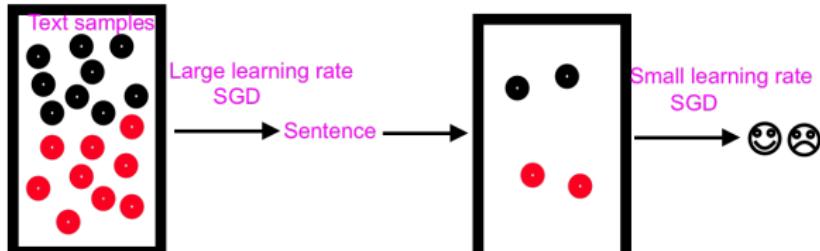


The intuition of SGD

- SGD with large learning rate further prevents memorization.


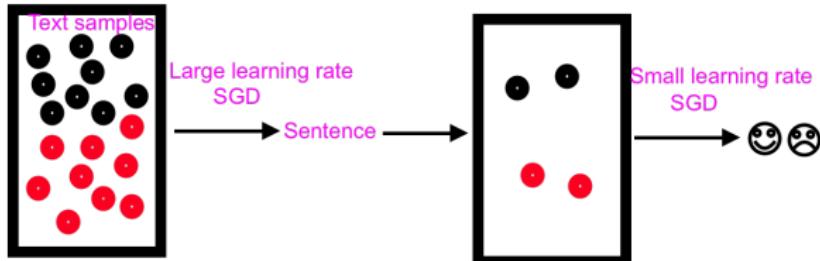
The intuition of SGD

- SGD with large learning rate further prevents memorization.



The intuition of SGD

- SGD with large learning rate further prevents memorization.



Memorization vs Generalization

- Memorization is not necessarily **bad!**

Memorization vs Generalization

- Memorization is not necessarily **bad**!
- I am Groot 😊 v.s. I am Groot 😈.

Memorization vs Generalization

- Memorization is not necessarily **bad!**
 - I am Groot ☺ v.s. I am Groot ☹.



Memorization vs Generalization

- Memorization is not necessarily **bad!**
 - I am Groot ☺ v.s. I am Groot ☹.



- The order matters: memorization after generalization.

The data set

- Data points in dimension $2p$: $x = (x_1, x_2)$ where $x_1, x_2 \in \mathbb{R}^p$.

The data set

- Data points in dimension $2p$: $x = (x_1, x_2)$ where $x_1, x_2 \in \mathbb{R}^p$.
- $x_1 \sim N(0, I)$ with probability $1 - \tau$, the label of x is $y(x) = 1_{\langle w, x \rangle \geq 0}$, where w is a vector in \mathbb{R}^p , otherwise $x_1 = 0$ and $y(x)$ is random (“sentence” signal).

The data set

- Data points in dimension $2p$: $x = (x_1, x_2)$ where $x_1, x_2 \in \mathbb{R}^p$.
- $x_1 \sim N(0, I)$ with probability $1 - \tau$, the label of x is $y(x) = 1_{\langle w, x \rangle \geq 0}$, where w is a vector in \mathbb{R}^p , otherwise $x_1 = 0$ and $y(x)$ is random (“sentence” signal).
- With probability $1 - \tau$ (τ is small),

$$x_2 = \begin{cases} z & \text{if } y(x) = 0; \\ z \pm \delta & \text{if } y(x) = 1. \end{cases}$$

(“special symbol” signal)

The data set

- Data points in dimension $2p$: $x = (x_1, x_2)$ where $x_1, x_2 \in \mathbb{R}^p$.
- $x_1 \sim N(0, I)$ with probability $1 - \tau$, the label of x is $y(x) = 1_{\langle w, x \rangle \geq 0}$, where w is a vector in \mathbb{R}^p , otherwise $x_1 = 0$ and $y(x)$ is random (“sentence” signal).
- With probability $1 - \tau$ (τ is small),

$$x_2 = \begin{cases} z & \text{if } y(x) = 0; \\ z \pm \delta & \text{if } y(x) = 1. \end{cases}$$

(“special symbol” signal)

- Remark: $\langle w, x \rangle$ can be replaced by two layer networks.

Our Theorem

Our Theorem

Theorem ((LM'19) Sketched)

*Given $N \approx p/\tau$ ($\tau \ll 1$) training examples, at training error (two layer over-parameterized neural network with **ReLU** activation) τ^2 for the cross entropy loss:*

Our Theorem

Theorem ((LM'19) Sketched)

Given $N \approx p/\tau$ ($\tau \ll 1$) training examples, at training error (two layer over-parameterized neural network with **ReLU** activation) τ^2 for the cross entropy loss:

- ① **SGD** with *large learning rate + learning rate decay* achieves generalization error $\leq \tau^{1+\Omega(1)}$.

Our Theorem

Theorem ((LM'19) Sketched)

Given $N \approx p/\tau$ ($\tau \ll 1$) training examples, at training error (two layer over-parameterized neural network with **ReLU** activation) τ^2 for the cross entropy loss:

- ① **SGD** with *large learning rate + learning rate decay* achieves generalization error $\leq \tau^{1+\Omega(1)}$.
- ② **SGD** with *small learning rate* has generalization error $\approx \tau$.

Our Theorem

Theorem ((LM'19) Sketched)

Given $N \approx p/\tau$ ($\tau \ll 1$) training examples, at training error (two layer over-parameterized neural network with **ReLU** activation) τ^2 for the cross entropy loss:

- ➊ **SGD** with *large learning rate + learning rate decay* achieves generalization error $\leq \tau^{1+\Omega(1)}$.
- ➋ **SGD** with *small learning rate* has generalization error $\approx \tau$.
- When $N \gg p$, they both **generalize** (as shown in the principle of over-parameterization), but **SGD** with *large learning rate + learning rate decay* **generalizes** even better.

Generalization:

```
def CMU_talk_example(x):  
    x = x^2  
    x = x + 10  
    x = x*3  
    x = x^(1/3)  
    x = x - 3  
    return x
```

Memorization:

```
def CMU_talk_example(x):  
    if x = 1, return 0.2  
    if x = 2, return 0.48  
    if x = 3, return 0.84
```

SGD with
small learning rate:
I like this one!

Generalization:

```
def CMU_talk_example(x):
    x = x^2.03
    x = x + 10.04
    x = x*3.02
    x = x^(1/3.01)
    x = x - 3.05
    return x
```

SGD with
large learning rate
(noisy)

almost correct

Memorization:

```
def CMU_talk_example(x):
    if x = 1.02 , return 0.2
    if x = 2.04,return 0.48
    if x = 3.01 , return 0.84
```

?????


x = 1

Principle of “noisy computation”

- Memorization in neural network: when $m \gg N$, by chance there will be N -neurons in the network that maps all data to **linear independent positions**. (Geometry of ReLU Lemma in [LL'18, ALS'18(a,b)])

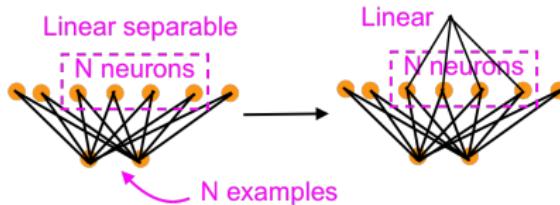
Principle of “noisy computation”

- Memorization in neural network: when $m \gg N$, by chance there will be N -neurons in the network that maps all data to **linear independent positions**. (Geometry of ReLU Lemma in [LL'18, ALS'18(a,b)])
- Then SGD just lazily finds a linear separator.

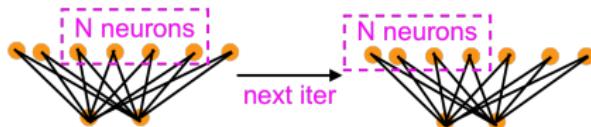


Principle of “noisy computation”

- Memorization in neural network: when $m \gg N$, by chance there will be N -neurons in the network that maps all data to **linear independent positions**. (Geometry of ReLU Lemma in [LL'18, ALS'18(a,b)])
- Then SGD just lazily finds a linear separator.

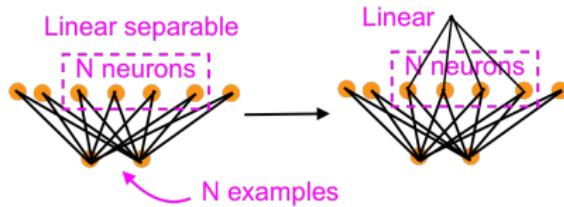


- SGD with large noise will make this set of N -neurons **changing rapidly**, and destroy the linear separator to prevent memorization.




Principle of “noisy computation”

- Memorization in neural network: when $m \gg N$, by chance there will be N -neurons in the network that maps all data to **linear independent positions**. (Geometry of ReLU Lemma in [LL'18, ALS'18(a,b)])
- Then SGD just lazily finds a linear separator.



- SGD with large noise will make this set of N -neurons **changing rapidly**, and destroy the linear separator to prevent memorization.



- This is what behind the principle of “noisy computation”.

Key Messages

- To the best of my knowledge, this is the **ONLY** work, that gives both efficiently achievable **lower and upper** bounds that large learning rate + learning rate decay improves **generalization** upon small learning rates.

Key Messages

- To the best of my knowledge, this is the **ONLY** work, that gives both efficiently achievable **lower and upper** bounds that large learning rate + learning rate decay improves **generalization** upon small learning rates.
- Even when the **training losses** of both methods are the same.

Key Messages


- To the best of my knowledge, this is the **ONLY** work, that gives both efficiently achievable **lower and upper** bounds that large learning rate + learning rate decay improves **generalization** upon small learning rates.
- Even when the **training losses** of both methods are the same.
- Prior works: small learning rate improves generalization. (More stable, lower training error etc.)

Key Messages

- To the best of my knowledge, this is the **ONLY** work, that gives both efficiently achievable **lower and upper** bounds that large learning rate + learning rate decay improves **generalization** upon small learning rates.
- Even when the **training losses** of both methods are the same.
- Prior works: small learning rate improves generalization. (More stable, lower training error etc.)
- This work: principle of “noisy computation”.

Key Messages

- To the best of my knowledge, this is the **ONLY** work, that gives both efficiently achievable **lower and upper** bounds that large learning rate + learning rate decay improves **generalization** upon small learning rates.
 - Even when the **training losses** of both methods are the same.
 - Prior works: small learning rate improves generalization. (More stable, lower training error etc.)
 - This work: principle of “noisy computation”.



Summary of my works

- I have described 10 percent of my works.

Summary of my works

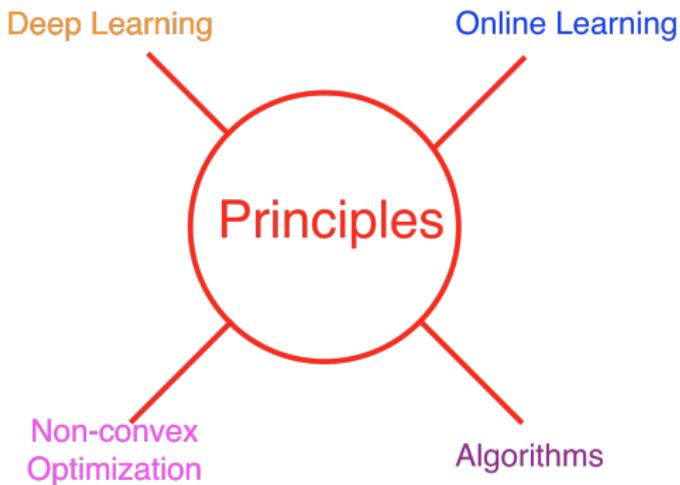
- I have described 10 percent of my works.
- But you might already see the rest from **Generalization**.

Summary of my works

- I have described 10 percent of my works.
- But you might already see the rest from **Generalization**.
- I work on the **principled** methods for machine learning. In the following four areas:

Summary of my works

- I have described 10 percent of my works.
- But you might already see the rest from **Generalization**.
- I work on the **principled** methods for machine learning. In the following four areas:



In each area

Deep Learning

Over-para: Training/Generalization

[LL'18, ALS'18, ALL'18]

Recurrent neural networks:

Training/Generalization [ALS'18, AL'18]

ResNet with/without over-para

[LY'17, AL'19]

Learning rate[LM'19]

In each area

Deep Learning

Over-para: Training/Generalization

[LL'18, ALS'18, ALL'18]

Recurrent neural networks:

Training/Generalization [ALS'18, AL'18]

ResNet with/without over-para

[LY'17, AL'19]

Learning rate[LM'19]

This talk

In each area

Deep Learning

- Over-para: Training/Generalization [LL'18, ALS'18, ALL'18]
- Recurrent neural networks:
Training/Generalization [ALS'18, AL'18]
- ResNet with/without over-para [LY'17, AL'19]
- Learning rate[LM'19]

Online Learning

- Convex bandit [HL'17]
- Chasing convex body[BLLS'18]
- First order/second order regret bounds for bandit/contextual bandit [ABL'18, BCL'18, BLLW'18]
- Distributed bandit[BLP'19]
- Bandit with approximation oracles [HHLL'18]

In each area

Deep Learning

- Over-para: Training/Generalization [LL'18, ALS'18, ALL'18]
- Recurrent neural networks: Training/Generalization [ALS'18, AL'18]
- ResNet with/without over-para [LY'17, AL'19]
- Learning rate[LM'19]

Online Learning

- Convex bandit [HL'17]
- Chasing convex body[BLLS'18]
- First order/second order regret bounds for bandit/contextual bandit [ABL'18, BCL'18, BLLW'18]
- Distributed bandit[BLP'19]
- Bandit with approximation oracles [HHLL'18]

Open problem since 2004

Open problem since 1991

Best student paper(ALT 18)

In each area

Deep Learning

- Over-para: Training/Generalization [LL'18, ALS'18, ALL'18]
- Recurrent neural networks: Training/Generalization [ALS'18, AL'18]
- ResNet with/without over-para [LY'17, AL'19]
- Learning rate[LM'19]

Online Learning

- Convex bandit [HL'17]
- Chasing convex body[BLLS'18]
- First order/second order regret bounds for bandit/contextual bandit [ABL'18, BCL'18, BLLW'18]
- Distributed bandit[BLP'19]
- Bandit with approximation oracles [HHLL'18]

Non-convex opt.

- Topic model/Matrix completion
- Matrix sensing/Tensor decomposition [LL'17, LLR'16(ab), LMZ'18, LZ'19]
- Mixture of regressions [LL'18]
- Graphical Models [LR'17]
- Escaping saddle points/approximate convex optimization[AL'18, LR'16]
- Experiments design[ALW'17]

In each area

Deep Learning

- Over-para: Training/Generalization
[LL'18, ALS'18, ALL'18]
- Recurrent neural networks:
Training/Generalization [ALS'18, AL'18]
- ResNet with/without over-para
[LY'17, AL'19]
- Learning rate[LM'19]

Online Learning

- Convex bandit [HL'17]
- Chasing convex body[BLLS'18]
- First order/second order regret bounds
for bandit/contextual bandit
[ABL'18, BCL'18, BLLW'18]
- Distributed bandit[BLP'19]
- Bandit with approximation oracles
[HHLL'18]

Non-convex opt.

- Topic model/Matrix completion
- Matrix sensing/Tensor decomposition
[LL'17, LLR'16, LLR'16, LMZ'18, LZ'19]
- Mixture of regressions [LL'18]
- Graphical Models [LR'17]
- Escaping saddle points/approximate
convex optimization[AL'18, LR'16]
- Experiments design[ALW'17]

With practical
improvements

Best paper(COLT 18)

In each area

Deep Learning

- Over-para: Training/Generalization
[LL'18, ALS'18, ALL'18]
- Recurrent neural networks:
Training/Generalization [ALS'18, AL'18]
- ResNet with/without over-para
[LY'17, AL'19]
- Learning rate[LM'19]

Online Learning

- Convex bandit [HL'17]
- Chasing convex body[BLLS'18]
- First order/second order regret bounds
for bandit/contextual bandit
[ABL'18, BCL'18, BLLW'18]
- Distributed bandit[BLP'19]
- Bandit with approximation oracles
[HHLL'18]

Non-convex opt.

- Topic model/Matrix completion
- Matrix sensing/Tensor decomposition
[LL'17, LLR'16, LLR'16, LMZ'18, LZ'19]
- Mixture of regressions [LL'18]
- Graphical Models [LR'17]
- Escaping saddle points/approximate
convex optimization[AL'18, LR'16]
- Experiments design[ALW'17]

Algorithms

- SVD(PCA), CCA, PCR, Online PCA
[AL'17(abcd)]
- Smoothed analysis[LS'18(ab)]
- Higher order acceleration/Distributed
Convex optimization[BJLLS'19(ab)]
- Matrix/Operator Scaling/Orbit
Closure Testing[ALOW'17,
AGLOW'18]
- L_p regression[BCLL'18]

In each area

Deep Learning

Over-para: Training/Generalization
[LL'18, ALS'18, ALL'18]
Recurrent neural networks:
Training/Generalization [ALS'18, AL'18]
ResNet with/without over-para
[LY'17, AL'19]
Learning rate [LM'19]

Online Learning

Convex bandit [HL'17]
Chasing convex body[BLLS'18]
First order/second order regret bounds
for bandit/contextual bandit
[ABL'18, BCL'18, BLLW'18]
Distributed bandit[BLP'19]
Bandit with approximation oracles
[HHLL'18]

One step towards derandomizing PT (complexity)

Non-convex opt.

Topic model/Matrix completion
Matrix sensing/Tensor decomposition
[LL'17, LLR'16, LLR'16, LMZ'18, LZ'19]
Mixture of regressions [LL'18]
Graphical Models [LR'17]
Escaping saddle points/approximate
convex optimization[AL'18, LR'16]
Experiments design[ALW'17]

Algorithms

SVD(PCA), CCA, PCR, Online PCA
[AL'17(abcd)]
Smoothed analysis[LS'18(ab)]
Higher order acceleration/Distributed
Convex optimization[BJLLS'19(ab)]
Matrix/Operator Scaling/Orbit
Closure Testing[ALOW'17,
AGLOW'18]
 L_p regression[BCLL'18]

Principles

Models

Over-parameterization

Algorithms

Inductive bias of SGD
Noisy computation

Principles

Models

Over-parameterization

Under-parameterization

Re-parameterization

Ensemble(Multiple Objectives)

Random initialization

Algorithms

Inductive bias of SGD


Noisy computation

Stable computation

Low dimension invariant

Symmetric breaking

Asynchronous



Principles

Models

Over-parameterization

Under-parameterization

Re-parameterization

Ensemble(Multiple Objectives)

Random initialization

First/Second order
bounds for contextual
/multi-arm bandit

Distributed bandit

Deep Learning
Learning rate

Algorithms

Inductive bias of SGD

Noisy computation

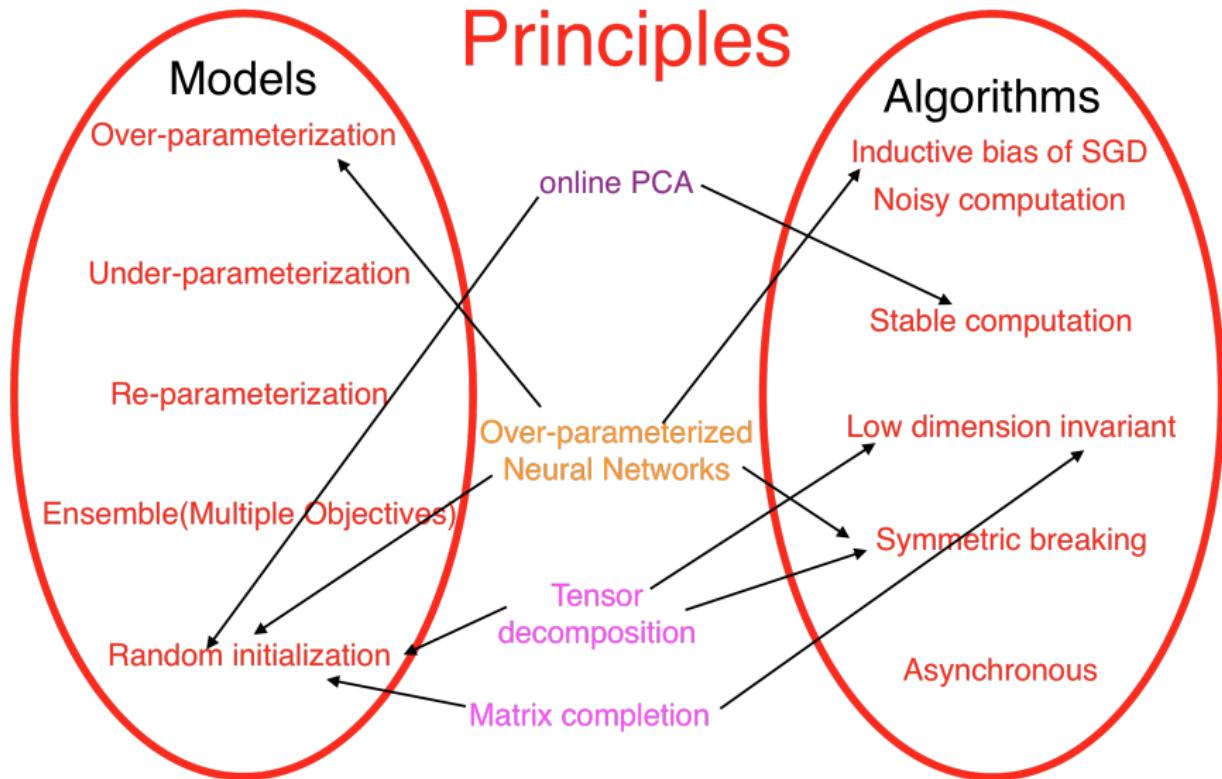
Stable computation

Low dimension invariant

Symmetric breaking

Asynchronous

Principles

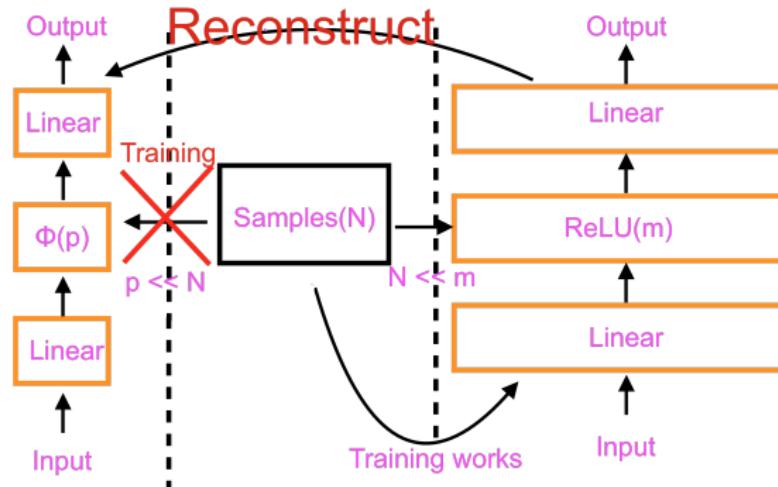


Future directions

- First understand the principles, then implement those principles more principally.

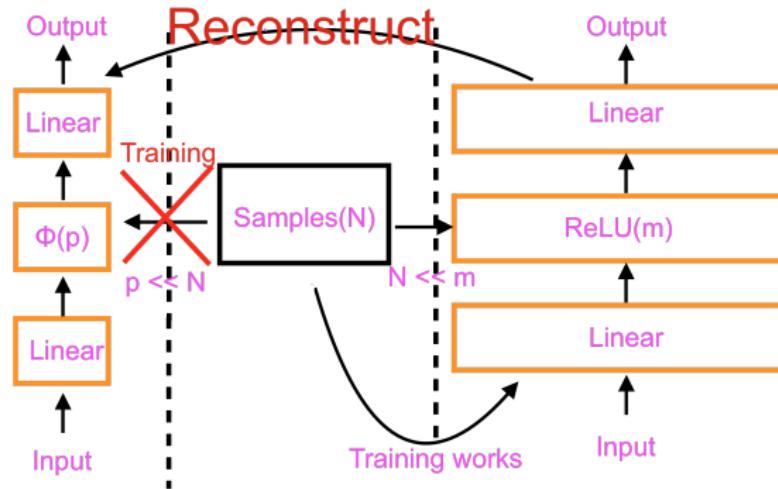
Future directions

- First understand the **principles**, then implement those **principles more principally**.
- Deep learning: reconstruct the small networks (Core features).



Future directions

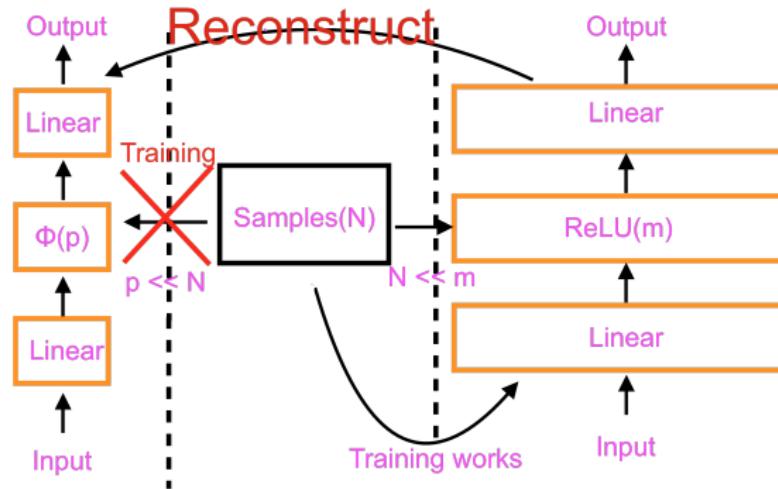
- First understand the **principles**, then implement those **principles more principally**.
- Deep learning: reconstruct the small networks (Core features).



- Deep learning: incorporate noise in a more controllable way. (SVRG + noise?)

Future directions

- First understand the **principles**, then implement those **principles more principally**.
- Deep learning: reconstruct the small networks (Core features).



- Deep learning: incorporate noise in a more controllable way. (SVRG + noise?)
- Much more exciting directions.



Yelp: Local Food & Services

Restaurant & Delivery Finder



4.1 ★★★★☆

4.42K Ratings

#10

Travel

12+

Age

What's New

Version 12.27.0

Version History

2d ago

We apologize to anyone who had problems with the app this week. We trained a neural net to eliminate all the bugs in the app and it deleted everything. We had to roll everything back. To be fair, we were 100% bug-free... briefly.